

Generative AI & LLM in Practice

Kishoreraj Jayakumar

The Nokia logo is displayed in white, uppercase letters within a dark blue circular area. This circle is part of a larger graphic consisting of two concentric white circles on a green-to-blue gradient background.

Self Introduction

KISHORERAJ JAYAKUMAR | Full Stack & Gen AI Engineer |



<https://www.linkedin.com/in/jkishoreraj/>

Agenda:

- Gen AI Introduction
- Gen AI Usage & Platforms
- LLM Overview
- Prompt Engineering
- Prompting Techniques
- RAG
- Agentic AI
- Short Demo



Beginner Questions

1. What is Gen AI?
2. How Gen AI works?
3. How to talk to an LLM?
4. How can custom knowledge be provided to LLM?
5. How to use LLM for Autonomous Use cases?

Generative AI in AI/ML

Artificial Intelligence

This is the overarching concept, aiming to create machines capable of performing tasks that typically require human intelligence, such as understanding language, recognizing images, and making decisions

Machine Learning

A subset of AI, ML focuses on algorithms that allow computers to learn from data without being explicitly programmed.

Deep Learning

A further specialization within ML, DL uses artificial neural networks with multiple layers (deep neural networks) to analyse data and identify complex patterns.

Generative AI

A type of AI that uses deep learning techniques to create new content, such as text, images, music, or video, based on learned patterns from existing data.

Generative AI Chatbots

some of the famous Gen AI Chatbots



ChatGPT

Owned By: **Open AI**

Models: GPT-3*, GPT-4*, GPT-O*



DALL·E

Owned By: **Open AI**

Models: DALL·E-2, DALL·E-3



Copilot

Owned By: **Microsoft**

Models: GPT-4*



GitHub Copilot

Dev By: **GitHub & Open AI**

Models: DALL·E-2, DALL·E-3



META AI

Owned By: **Meta**

Models: Llama 3.2, Llama 3.1 -405B



Claude

Owned By: **Anthropic**

Models: Haiku, Sonnet, Opus



Owned By: **Google**

Models: Gemini -1.5 & 2.0
(Flash/Pro/Lite)

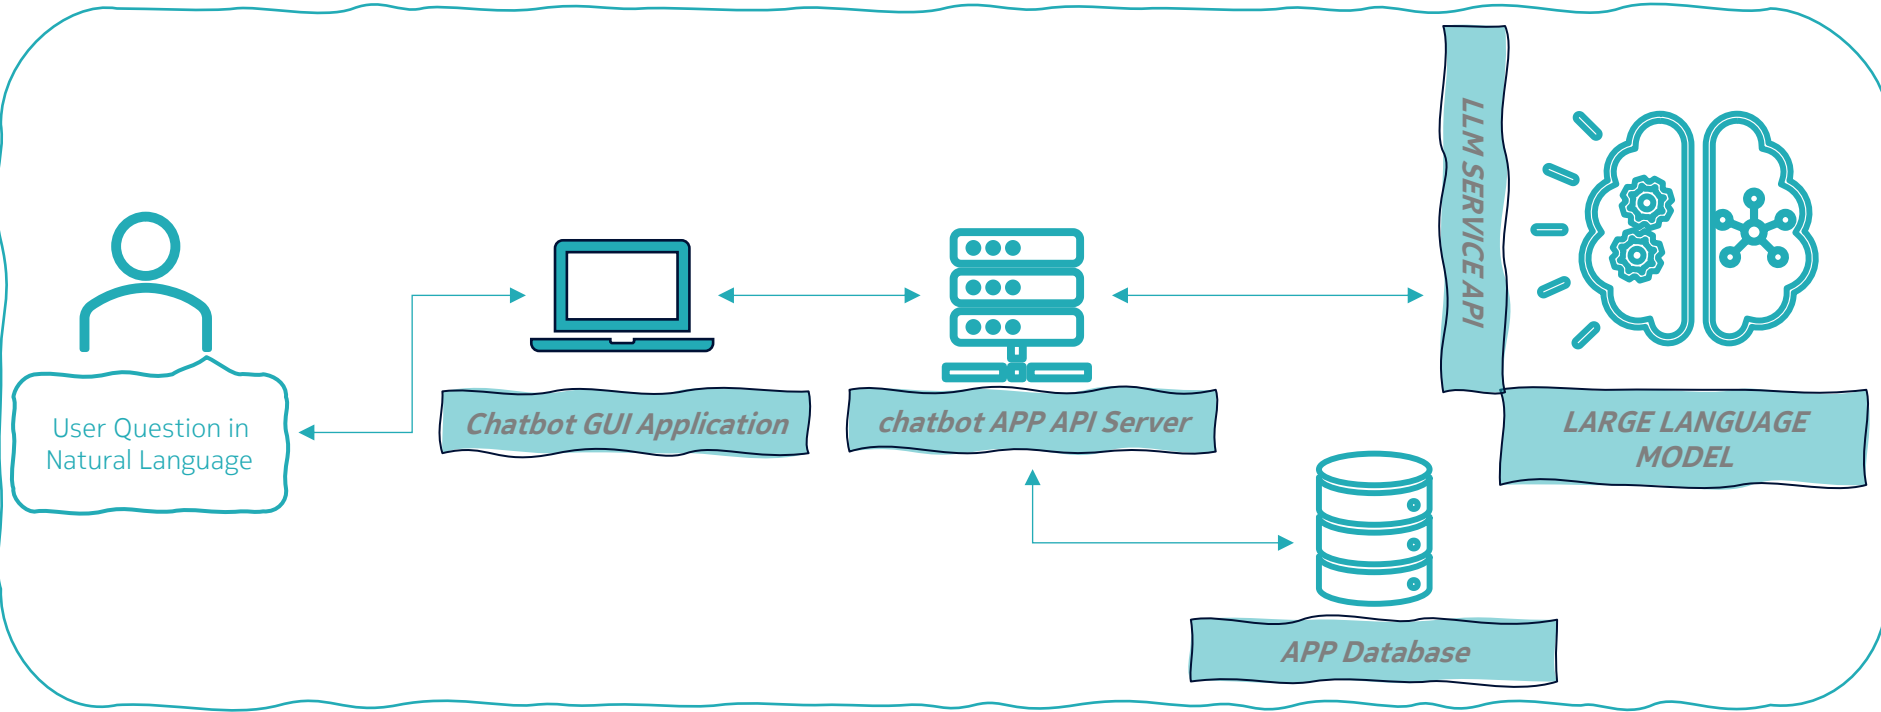


perplexity

Owned By: **Perplexity AI**

Models: sonar pro/reasoning

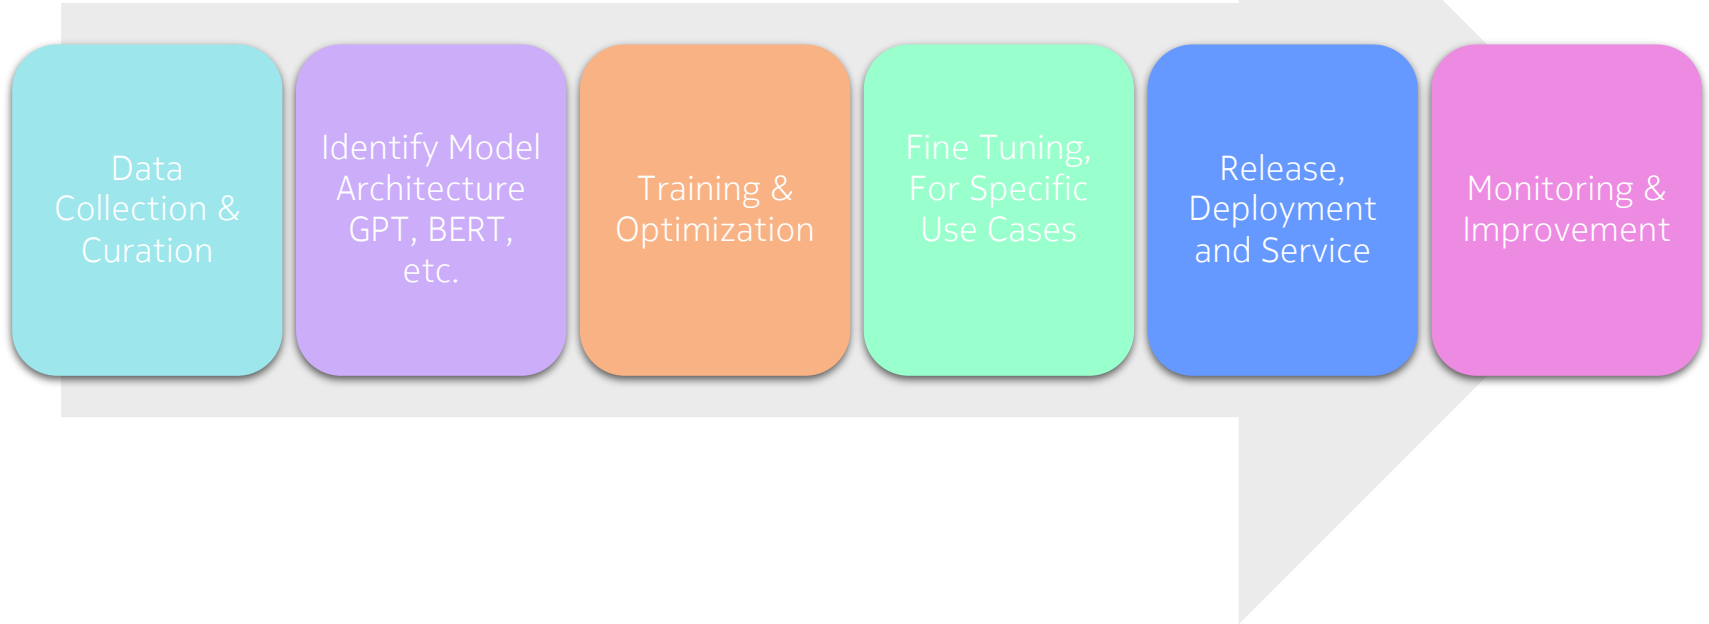
Gen-AI App – High-Level Design



Large Language Models (LLMs)

- AI/ML Model built with deep learning techniques
- Neural network model trained with massive data
- Understands, interprets, and generates human-like text
- Mostly uses transformers architecture, and undergoes curation, and tokenization processes while training
- Knowledge level is commonly identified by “parameters”,
 Ex: Llama 8B , gemma 27B
- Widely used for text generation, language translations, chatbots, summarization and data reasoning, etc.,
- Also, multimodal LLMs are used for text, image, audio, video generation and interpretations

LLM Development Overview



Open Vs Closed source LLMs

Aspect	Open-Source LLMs 🛠️	Closed-Source LLMs 🔒
Access	Free to use, modify, and distribute	Restricted access, proprietary licenses
Transparency	Fully transparent, code and training data available	Opaque, internal algorithms and data not disclosed
Customization	Fully customizable for specific needs	Limited customization (depends on API features)
Deployment	Can be deployed on-premise or in a private cloud with GPUs by ourselves	Hosted by provider, SaaS-based API access
Data Privacy	Full control over data and model usage	Data may be logged or used for improvement
Security	User-controlled security measures	Provider ensures security, but data exposure risks exist
Cost	Free to use, but infrastructure costs apply	Typically pay-per-use or subscription-based
Performance	Varies based on hardware and optimization	Generally well-optimized and high-performing
Scalability	Requires own infrastructure for scaling	Scales easily with provider's cloud infrastructure
Community Support	Strong open-source communities (Hugging Face, Meta, Mistral)	Limited, mostly official support
Regulatory Compliance	Easier to ensure compliance by self-hosting	Compliance depends on provider's policies
Updates & Maintenance	Requires manual updates and improvements	Automatically updated by provider
Examples	LLaMA, Mistral, Falcon, GPT-NeoX, BLOOM	GPT-4, Gemini, Claude, CoPilot

LLM Integration - Key Techniques:



Prompt Engineering

Crafting
Effective Inputs



RAG


Retrieval Augmented Generation

Providing
External Data



Fine-Tuning

Customizing a
Pre-trained LLM



Agentic AI

Enabling LLM
response to
trigger actions


Prompt

A prompt is a query, command, or statement guiding the LLM to generate a desired output

The LLM model analyzes the prompt and uses its training data to produce a response that is relevant and appropriate to the prompt.

Prompt Engineering

1. The art and science of crafting input prompts to maximize the quality of output from language models.
2. Sets the context and intention
3. Directly influences model accuracy and relevance.
4. It improves communication between humans and machines, ensuring the resulting interaction is efficient and effective.



Why Does Prompt Engineering Matter?



Precision: Reduces ambiguity in model output.



Efficiency: Minimizes trial-and-error cycles.



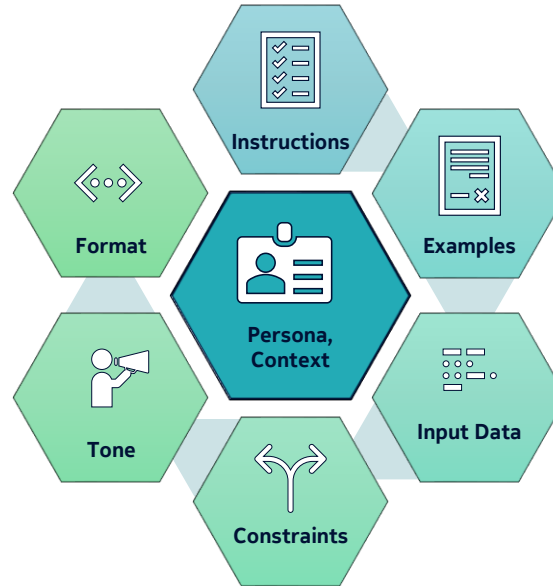
Adaptability: Tailors responses to specific domains.



Example: "Summarize this text" vs. "Summarize this text in 3 bullet points for a business audience."

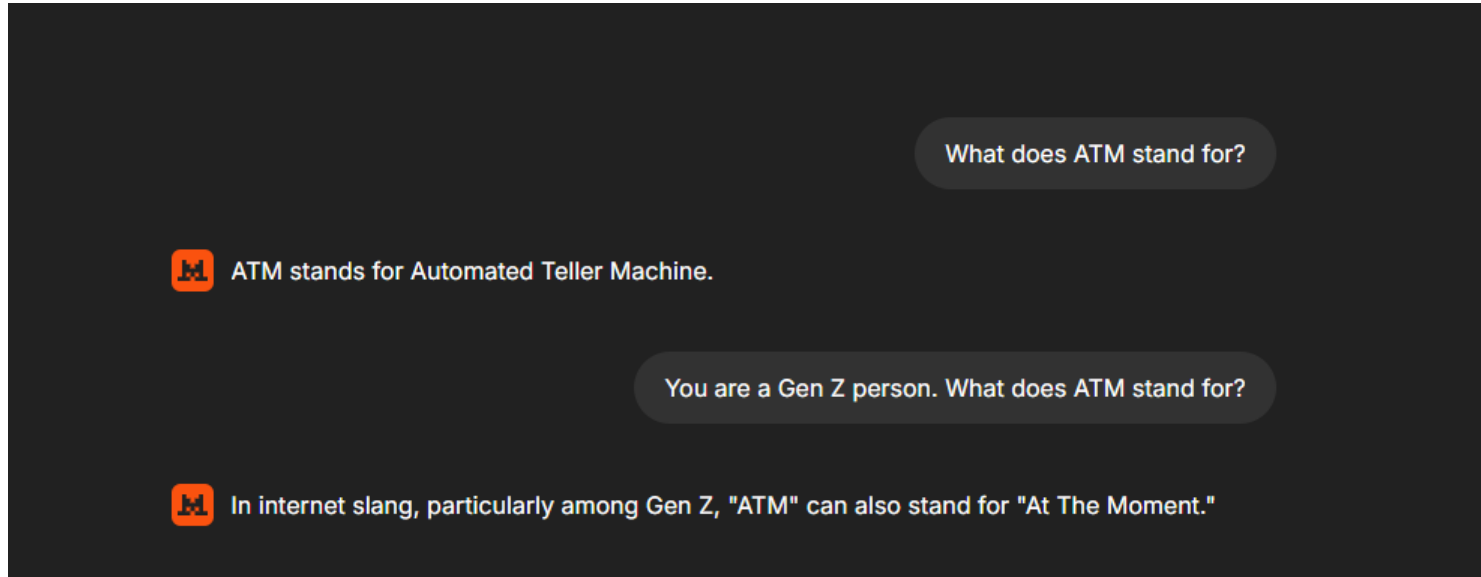
Prompting Elements

Different sections of a prompt, helps in guiding LLM for a desired response



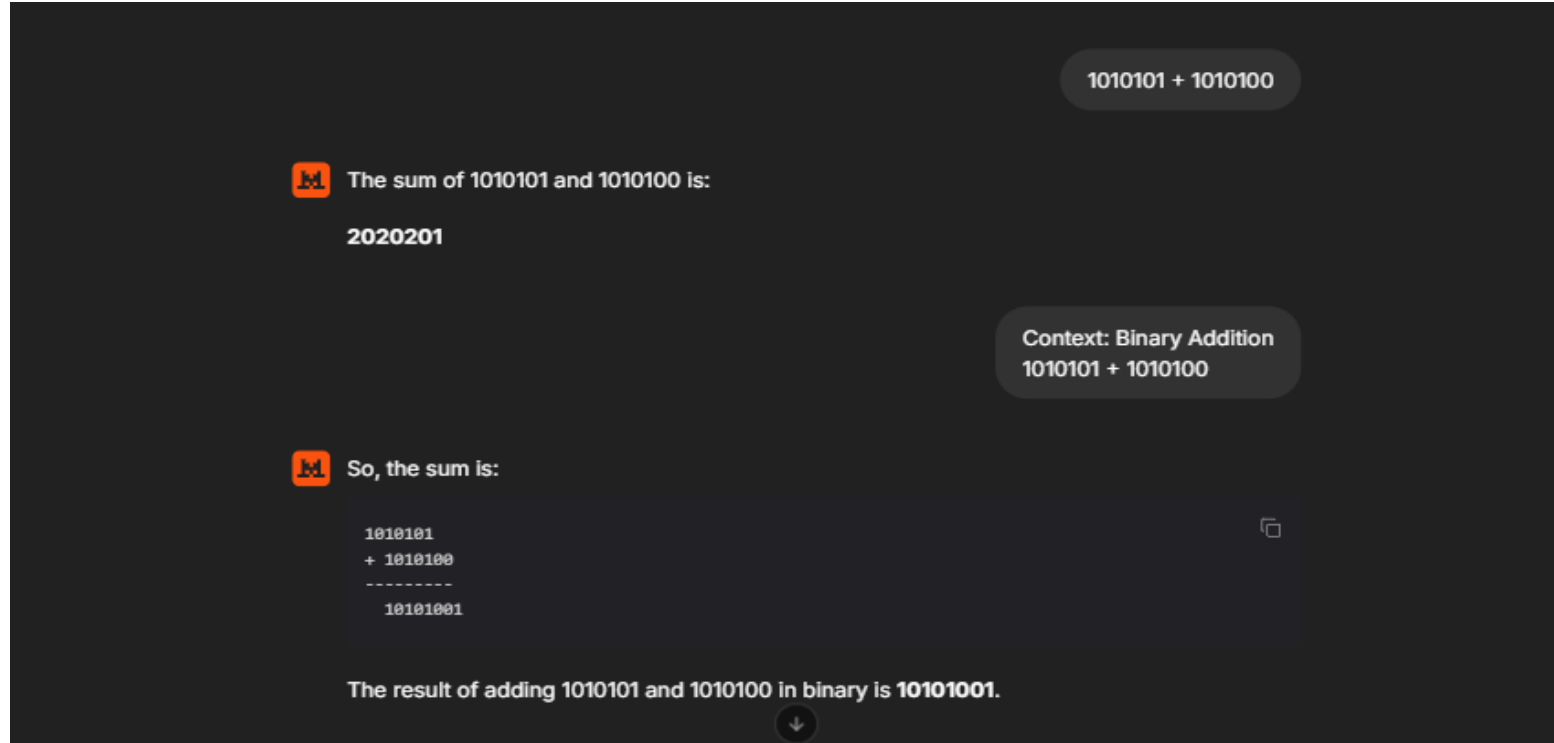
Prompt Structure

Role



Prompt Structure

Context



The screenshot displays a chat interface with a dark background. At the top right, a grey bubble contains the expression $1010101 + 1010100$. Below this, a message from a user (indicated by an orange icon) asks for the sum of 1010101 and 1010100. The assistant's response (indicated by a blue icon) provides the decimal sum 2020201. A second grey bubble contains the context "Context: Binary Addition" and the same expression. The user then asks for the sum in binary. The assistant responds with a binary addition diagram showing the sum of 1010101 and 1010100 as 10101001, followed by the text "The result of adding 1010101 and 1010100 in binary is 10101001." and a downward arrow icon.

$1010101 + 1010100$

M The sum of 1010101 and 1010100 is:

2020201

C Context: Binary Addition
 $1010101 + 1010100$

M So, the sum is:

```
1010101
+ 1010100
-----
10101001
```

The result of adding 1010101 and 1010100 in binary is **10101001**.

Prompt Structure

Final Example

The screenshot shows a chat interface with a dark background. At the top, a light gray box contains the following prompt structure:

- context:** Summarise the news article
- input :** "US President Donald Trump on Sunday announced that he will impose a 25 per cent tariff on all steel and aluminium imports, with details expected on Monday. Speaking to reporters aboard Air Force One on Sunday, Trump confirmed that the new tariffs would apply to all countries but did not specify when they would take effect, reports news agency Bloomberg. The tariff announcement comes just before Prime Minister Narendra Modi's scheduled visit to the United States on February 12, adding a layer of complexity to trade discussions between the two nations. The president also stated that he would unveil reciprocal tariffs on nations that impose taxes on US imports later in the week. These additional tariffs are expected to follow shortly after their announcement, potentially by midweek."
- output :** provide in json format {
 summary : <summary>
 word_count : <word_count>
}
- restriction :** summary should not exceed length of 100 characters
- scope:** first take key points, summarise with restricted length and provide word count

Below the prompt box, there is a chat history section. It shows a message from a user (indicated by an 'M' icon) with the following JSON output:

```
{  
  "summary": "Trump announces 25% tariff on steel and aluminium imports, effective date unspecified",  
  "word_count": 20  
}
```

The chat interface includes standard icons for editing, deleting, and copying messages, as well as a timestamp of 18:47.

Prompting Techniques

Different ways of crafting prompts based on requirements

Zero Shot

Direct questions without any examples

Example: "Explain black hole in simple terms"

Few Shot

Provide a few examples before asking the model to generate a response

Example: "*"Translate the following: 1. Hello → Hola 2.Thank you → Gracias 3.Good night → ???"*

Response: "Buenas noches."

Role Based

Assigning a specific role to the AI for contextual responses.

Example: "You are a South Indian grandmother. Explain why filter coffee is the best."

Response: "Ayyo! Nothing beats fresh decoction and thick milk—better than any machine coffee!"

Chain of Thought

Encourages step-by-step reasoning before answering.

Example: "If a train moves at 50 km/h for 3 hours, how far does it go? Think step by step."

Response: "Step 1: Speed = 50 km/h - Step 2: Time = 3 hours Step 3: Distance = Speed × Time = 50 × 3 = 150 km."

Self Consistent – CoT

The model generates multiple CoT answers for n iterations of CoT prompts and picks the most consistent one

Use Case: Mathematical problems, logic puzzles, and complex decision-making.


Tree of Thought

Instead of a single CoT path, the model branches into multiple reasoning paths, evaluates them, and picks the best.


Example: "How can I prepare for my college final exams effectively? Think in multiple ways before deciding the best approach."

Response: "Path 1: Structured Study Plan, Path 2: Group Study Approach, Path 3: Visual & Memory Techniques150 km."


Prompting Illustrations




(a) Input-Output Prompting (IO)



(c) Chain of Thought Prompting (CoT)



(c) Self Consistency with CoT (CoT-SC)



(d) Tree of Thoughts (ToT)

Prompting – Output Controlling Attributes

Parameter	Definition	Effect	Value Range	Use Cases
Temperature 🌡️	Controls randomness in responses.	Lower values (0.1-0.5) make responses deterministic , higher values (0.7-1.5) increase creativity .	0.0 to 2.0	Code generation (low), Creative writing (high)
Top-k Sampling 🎯	Limits word selection to the k most probable words.	Lower k makes responses focused , higher k makes them diverse .	1 to 100	Factual answers (low k), Storytelling (high k)
Top-p Sampling 🏆	Chooses words based on cumulative probability.	Lower p makes responses conservative , higher p allows more diverse outputs.	0.0 to 1.0	Formal writing (low p), Casual conversations (high p)
Frequency Penalty 📄	Reduces repetition of frequently used words.	Higher values discourage repeated phrases.	0.0 to 2.0	Prevents repetitive responses in essays, articles
Max New Tokens 📄	Limits the number of generated tokens (words).	Higher values allow longer responses , lower values keep them concise .	10 to 4096+	Short summaries (low), Long-form writing (high)

Prompting Demo






Retrieval Augmented Generation (RAG)

NOKIA

Retrieval Augmented Generation (RAG)



What?

RAG is an AI technique that enhances text generation by retrieving relevant external knowledge before generating responses

Why?

Traditional LLMs rely only on pre-trained knowledge, which can become outdated. RAG **dynamically retrieves up-to-date information**, ensuring more factual and context-aware responses.

When?

It is used when AI needs to **answer domain-specific, dynamic, or real-time queries**, such as **legal, medical, or technical** support

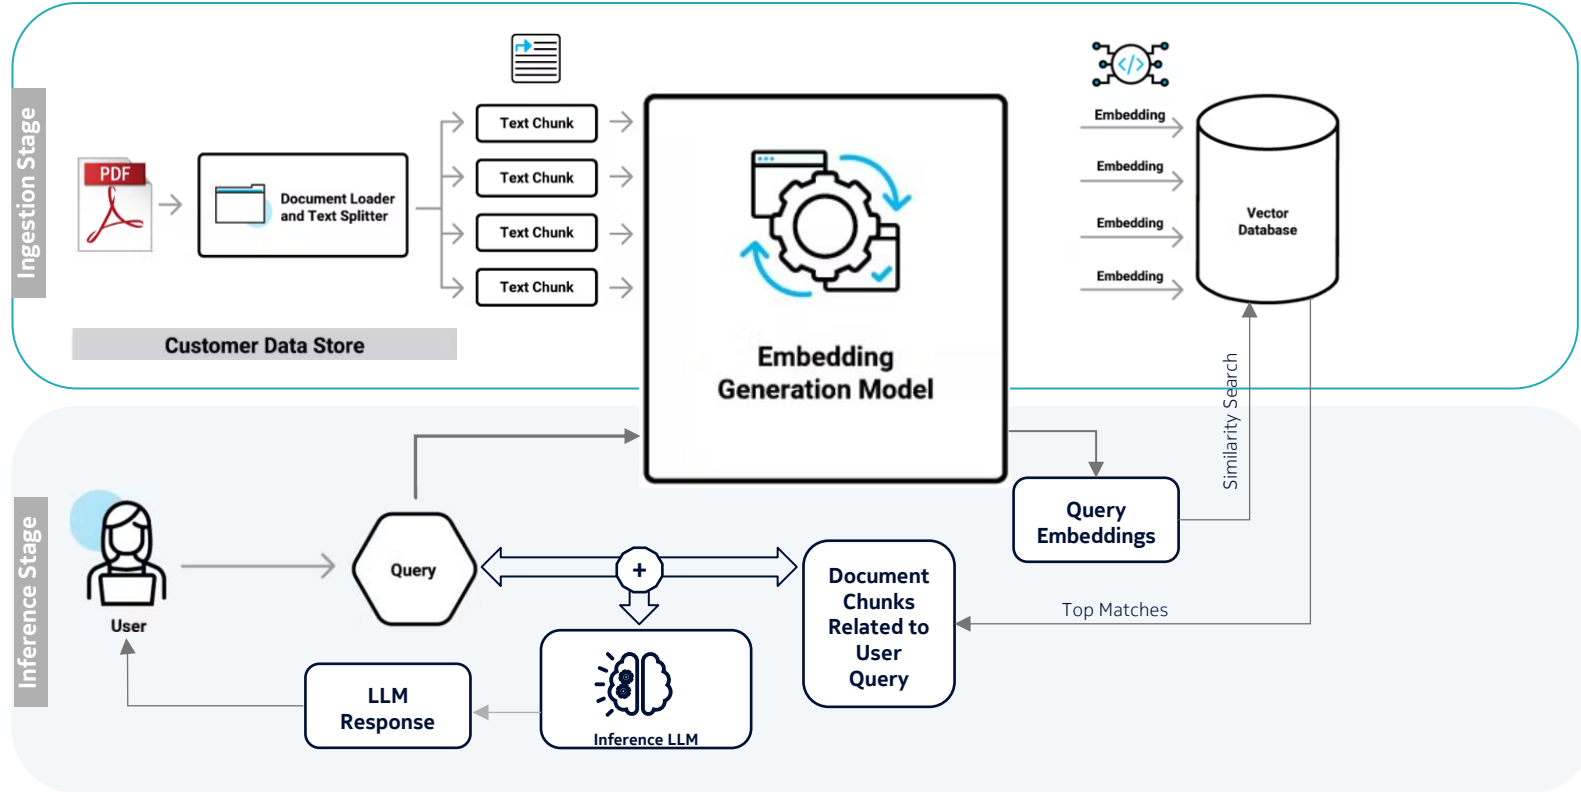
Where?

Commonly applied in **customer service chatbots, enterprise knowledge retrieval, AI search engines, and research assistants** that require factual grounding.

How?

RAG first **retrieves** the most relevant data from a **vector database or document store** and then **uses an LLM** to generate an informed response, improving accuracy and reducing hallucinations.

RAG Implementation




```
mirror_mod.mirror_object = mirror_mod.mirror_object
operation == "MIRROR_X":
    mirror_mod.use_x = True
    mirror_mod.use_y = False
    mirror_mod.use_z = False
operation == "MIRROR_Y":
    mirror_mod.use_x = False
    mirror_mod.use_y = True
    mirror_mod.use_z = False
operation == "MIRROR_Z":
    mirror_mod.use_x = False
    mirror_mod.use_y = False
    mirror_mod.use_z = True
```

```
#selection at the end -add
mirror_ob.select= 1
modifier_ob.select=1
context.scene.objects.active = mirror_ob
("Selected" + str(modifier_ob.name))
mirror_ob.select = 0
bpy.context.selected_objects = [mirror_ob]
data.objects[one.name].select = 1
```

DEMONSTRATION

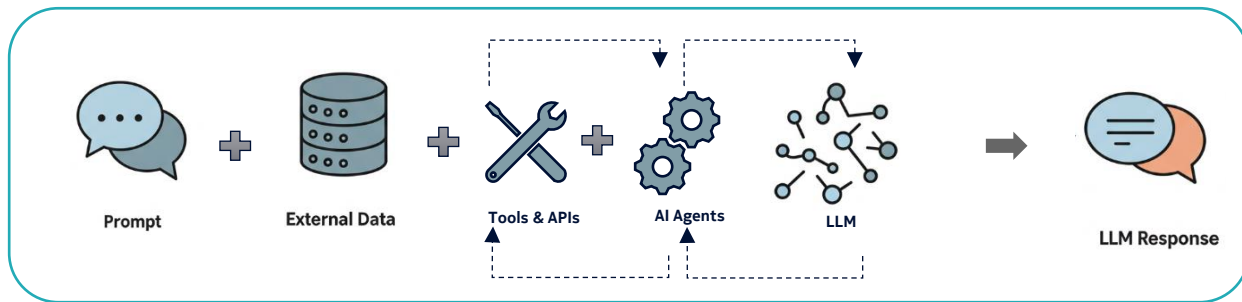
```
types.Operator):
    X mirror to the selected
    object.mirror_mirror_x"
    mirror X"
```

A man in a light blue shirt is leaning over a desk, smiling and looking at a humanoid robot. The robot is white with black accents and is sitting at the desk, looking up at the man. The man's hands are on a keyboard. The background is a blurred office setting with plants and computer monitors.

AGENTIC AI

Agentic AI

Extending LLM capabilities with AI Agents



- Agentic AI is a software system designed to interact with data and tools in a way that requires minimal human intervention.
- Agentic AI is sometimes referred to as autonomous AI. This is because it has the capability to communicate and collaborate with other AI systems and digital infrastructures on behalf of a human user or another AI agent
- To bring agentic AI to practice, you create a system that provides an LLM with access to external tools and algorithms that supply instructions for how the AI agents should use those tools

Agentic AI – System Design



Agent

A software program that uses artificial intelligence (AI) to perform tasks on behalf of a user



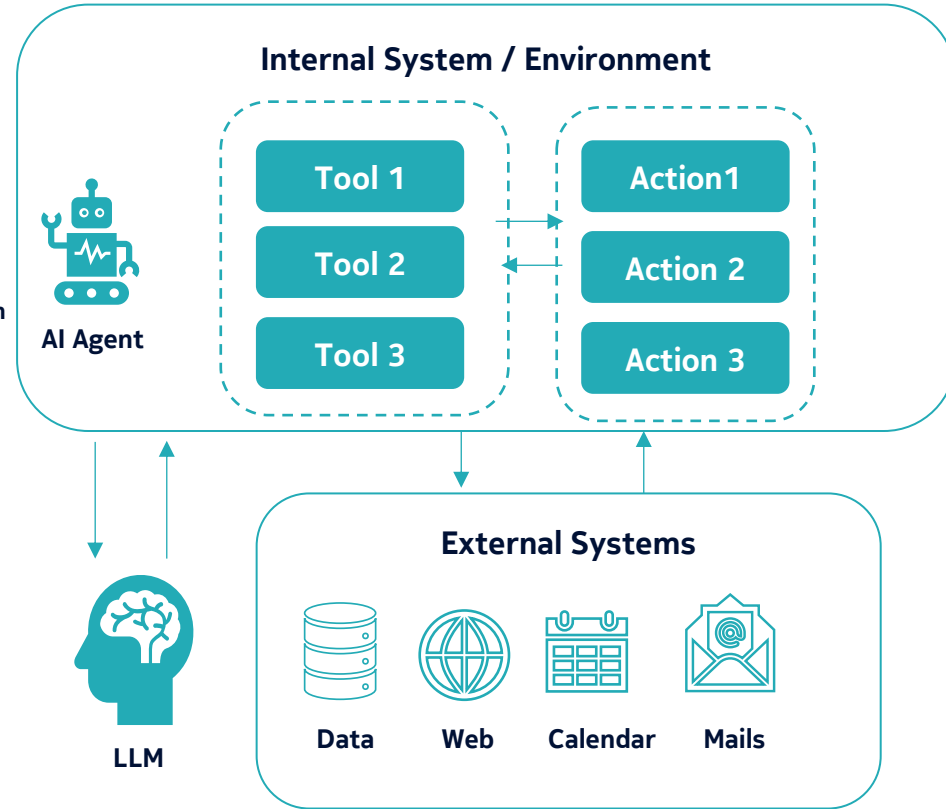

Tools

A "Tool" refer to functions or APIs that enable the agent to interact with the external world or another software program



Task

A "task" refers to a specific, predefined goal or objective that an AI agent is programmed to achieve through interacting with its environment



Agentic AI – Real Life Examples

1. **Travel Planner** – Plans itinerary, logistics, accommodations
2. **Customer Service Bot** – Handles support queries and processes refunds.
3. **AI Research Assistants** – Fetches and summarizes data for users.
4. **Automated Coding Agents** – Writes, tests, and debugs code.
5. **Smart Assistants (like Jarvis)** – Plans and executes tasks proactively.



Q & A

Thank You

Endorsements

<https://www.linkedin.com/in/jkishoreraj/>

Kishoreraj Jayakumar
Technical Specialist - Nokia | Full Stack
Engineer | Gen AI Enthusiast



Survey

<https://forms.office.com/e/AkyCePZqbU>

